

CHARLES DARWIN UNIVERSITY

Assessment 1 - Group Project Proposal and Design Plan

Group Number – SYD7

Student Name	Student ID
Guvvala Thejeswar Reddy	384042
Julakanti Thirumala Raja Sumanth Reddy	386402
Dhruv Pramod Vakharia	385234

Assessment 1: Project Proposal and Design (10%)

Field	Detail
Project Title	Flight Delay Prediction: A Predictive Analytics and Forecasting Approach
Project Theme	Theme 2 – Predictive Analytics and Forecasting
Unit Code	PRT661
Lecturer	Reem Sherif
Domain	Aviation Operations & Transport Analytics
Primary Dataset	U.S. DOT BTS Airline On-Time Performance Dataset (2018–2023)
Tools	Python, DuckDB, MinIO, scikit-learn, XGBoost, LightGBM, Plotly Dash, Apache Airflow

Table of Contents

<u>1. PROJECT OVERVIEW AND OBJECTIVES</u>	<u>5</u>
<u>2. THEME JUSTIFICATION</u>	<u>6</u>
<u>3. PROBLEM STATEMENT</u>	<u>6</u>
3.1 TARGET USERS AND BUSINESS CASE.....	6
<u>4. INITIAL ARCHITECTURE DIAGRAM</u>	<u>7</u>
<u>5. WORKFLOW PLAN</u>	<u>7</u>
<u>6. DATASET SELECTION AND DATA CHARACTERISTICS</u>	<u>8</u>
6.1 PRIMARY DATASET	8
6.2 FIVE VS ANALYSIS	8
<u>7. TASK ALLOCATION AND TEAM STRUCTURE</u>	<u>9</u>
<u>8. INITIAL RISK ANALYSIS</u>	<u>10</u>
<u>9. ETHICS, PRIVACY, AND SECURITY CONSIDERATIONS</u>	<u>10</u>
9.1 ETHICS AND FAIRNESS.....	10
9.2 PRIVACY	10
9.3 SECURITY.....	11
<u>10. REFERENCES</u>	<u>12</u>
<u>APPENDICES</u>	<u>13</u>
APPENDIX A: DATA DICTIONARY	13

APPENDIX B: DETAILED PIPELINE PROCESS DESIGN 14

APPENDIX C: MODEL DEVELOPMENT AND EVALUATION PLAN 15

APPENDIX D: TECHNOLOGY JUSTIFICATION MATRIX..... 15

1. Project Overview and Objectives

There is a significant economic cost associated with flight delays. The societal benefits of a 10% decrease in U.S. flight delays would amount to USD 17.6 billion, and those benefits would be USD 38.5 billion with a reduction of 30%, after taking into consideration airline costs, passenger time, and lost productivity, find Peterson et al. (2013). This project involves the development of a batch predictive analytics system that classifies whether a scheduled U.S. domestic flight is expecting a delay and estimates the duration of the expected delay based on historical flight data that is enhanced with weather, carrier and airport context data, processed in a five-step pipeline and presented to airline operations teams on an interactive dashboard.

Lambelho et al. (2020) present evidence of machine-learning models outperforming simple scheduling heuristics for predicting delay/cancellation with carrier behaviour and contextual features; and that gradient-boosting ensembles (XGBoost, LightGBM, CatBoost) with weather features can achieve strong accuracy on large-scale aviation datasets, as demonstrated in Dai (2024). This project brings these results into practice in the form of a reproducible, open-source pipeline.

Objectives: acquire, clean and merge BTS flight records, NOAA weather and airport metadata (2018–2023); perform EDA on BTS flight delay patterns across carriers, airports, routes, seasons, and NOAA weather; prepare temporal and contextual predictive features; train and test classifiers (Logistic Regression, Random Forest, XGBoost, LightGBM) for binary prediction delay or not; create a regression sub-model forecasting duration of delay in minutes; provide the results through an operational dashboard; and apply ethical, privacy-aware, and reproducible practice throughout.

This is a Data Science endeavour, not just a Data wrangling/ML exercise, everything from acquiring and storing data, to wrangling, exploring, modelling, to decision-oriented visualization. The decision that is supported is well-specified and measurable: will this flight depart on time, and if not, about how many minutes late will it be?

2. Theme Justification

This project falls into Theme 2 – Predictive Analytics and Forecasting: batch, structured, tabular data for supervised learning. The following six, mandated but often overlooked, components of a pipeline acquisition, storage, processing, analytics/ML, visualisation and workflow automation are all addressed as well as data-quality assessment, data Five Vs analysis, and dashboard design for stakeholders.

3. Problem Statement

Real world operations rely on rules-of-thumb based aviation delay prediction that assumes a stochastic weather process, fails to take network interactions between flights into account and is based on an aviation sector information system with considerable heterogeneity across carriers and airports. The research problem is as follows: Can a machine-learning system predict, before departure, whether a flight will be delayed and, if it is, when estimate the duration of the delay based on historical flight performance information that includes the weather at the time of flight, the carrier information, and information about the airports where the flight takes off and lands?

The flight delay patterns in the COVID-19 period are materially different from those of the previous years, Faiza and Khalil (2023) demonstrate, leading to temporal validation: this project employs walk-forward cross validation and also does not consider the data for the COVID period (2020–2021) with other normal years. The authors of Li et al. (2024) conclude that the topology of airports, namely the connectivity of a hub and its node centrality are measurable influences for the propagation of delays and thus their graph-derived features, such as node centrality and hub classification, are motivated for use in the feature set.

3.1 Target Users and Business Case

While the primary users will be the airline operations control centres and the airport slot coordinators, the secondary users will be the operations, passenger facing apps and the airline revenue management. Even for small delays of a few minutes on a system-wide level, the benefits in recuperated welfare are measured in billions, while on a carrier level

the predictive delay intelligence makes pro-active crew reallocation and gate reallocation, and passenger rebooking, possible.

4. Initial Architecture Diagram

It is built around a batch analytics pipeline in five open-source layers (seen in Figure 1).




Figure 1: Five-Layer Open-Source Batch Analytics Architecture

The data flow is: BTS/NOAA sources to Python ingestion to MinIO raw zone to DuckDB processing to Parquet feature store to ML pipeline to Plotly Dash dashboard. The full history of all transformations is recorded in a data lineage document and uploaded onto GitHub, based on the work and suggestions of Stoudt et al. (2024) about an auditable and responsible data science workflow.

5. Workflow Plan

The project is broken down into four sprints that correspond to the assessment calendar, and planning, review, and retrospective ceremonies are conducted every two weeks and recorded in Microsoft Teams and tracked on Jira.

Table 1: Sprint workflow aligned with PRT661 assessment calendar.

Sprint	Weeks	Key Activities	Deliverable
1	1–3	Dataset audit; repo/Jira setup; EDA kick-off; initial architecture	Assessment 1 Proposal
2	4–7	Cleaning pipeline; feature engineering; baseline models; Airflow DAG; MinIO config	Assessment 2 Progress Report

3	8–11	XGBoost/LightGBM/RF; Optuna HPO; SHAP; fairness check; SMOTE; dashboard prototype	Assessment Presentation	3
4	12–14	Dashboard finalisation; automated reporting; ethics review; lineage docs; README	Assessment Final Report	4

6. Dataset Selection and Data Characteristics

6.1 Primary Dataset

The key data source is the U.S. DOT Bureau of Transportation Statistics (BTS) Airline On-Time Performance Dataset ([transtats.bts.gov](https://www.transtats.bts.gov)) that includes scheduled/actual times, delay minutes by cause, origin/destination, tail numbers, and distance for air transportation within the United States. The extract from 2018–2023 has about 36 million records in 18 carriers and 373 airports (~12 GB CSV; ~2.5 GB in Parquet format). Weather and geographic/hub additional data are provided by NOAA GHCN-Daily weather and the OpenFlights airport database.

Primary Dataset:

BTS Airline On-Time Performance Dataset

<https://www.transtats.bts.gov>

Supplementary

NOAA GHCN-Daily weather

<https://www.ncei.noaa.gov/products/land-based-station/global-historical-climatology-network-daily>

OpenFlights airport metadata

<https://openflights.org/data.php>

6.2 Five Vs Analysis

Table 2: Five Vs analysis of the flight delay prediction dataset.

Dimension	Assessment and Architectural Implication
Volume	~36M records / ~12 GB CSV to ~2.5 GB Parquet; DuckDB in-process SQL; MinIO three-zone storage.
Velocity	Batch: BTS monthly, NOAA daily via API. No streaming needed; Airflow manages scheduling.
Variety	Structured CSV, semi-structured JSON (NOAA), relational lookups (OpenFlights) unified via DuckDB.
Veracity	High — mandated federal reporting. ~1.3% missing delay-cause codes; 2020–2021 needs stratification.
Value	Supports delay classification, duration regression, carrier benchmarking, seasonal/airport monitoring.

7. Task Allocation and Team Structure

This is because the team adopts a distributed ownership model: each member owns a layer of the pipeline, and is responsible for designing, implementing, documenting (using GitHub/Jira tickets, see Table 3), and testing this layer. Sprint ceremonies are conducted biweekly on Microsoft Teams and each portion of work is backed by individual commit, pulls request and or jira tasks for contribution evidence.

Table 3: Team roles and deliverables.

Member	Role	Deliverables
1	Data Engineer / Pipeline Lead	BTS + NOAA ingestion; MinIO config; Airflow DAG; ETL; data quality audit
2	Data Analyst / EDA Lead	EDA notebooks; delay visualisations; DuckDB SQL analytics
3	ML Engineer	Model training; Optuna HPO; walk-forward CV; SHAP; fairness evaluation
4	Dashboard Developer	Plotly Dash app; prediction UI; performance views; scheduled reports
5	Project Manager /	Jira management; architecture diagrams; ethics

	Ethics Lead	review; README/risk register
--	-------------	------------------------------

8. Initial Risk Analysis

Table 4: Risk register with likelihood, impact, and mitigation strategies.

Risk	Likelihood	Impact	Mitigation
Dataset size exceeds local memory	High	Medium	Parquet + DuckDB; prototype on 2022–2023 subset
COVID-19 period causes distribution shift	High	High	Temporal stratification; walk-forward CV
Class imbalance (~80/20) biases classifier	Certain	High	SMOTE/class weighting; evaluate via AUC-ROC/F1
NOAA API rate limits or outages	Medium	Medium	Cache responses; fallback to NOAA bulk FTP files
Team member non-participation	Medium	High	GitHub commit evidence; redistribute tasks per sprint
Scope creep into real-time streaming	Medium	Medium	Strict backlog governance; Jira-documented scope changes

9. Ethics, Privacy, and Security Considerations

9.1 Ethics and Fairness

The BTS data does not include any PII for passengers, but does provide data on the carriers and routes that could show disparities influenced by socioeconomic access to air travel. Drawing on Mehrabi et al. (2021) fairness framework, model performance will be broken down based on carrier type, airport hub classification, and geography of the route, and gaps reported transparently, with post hoc interventions, ranging from equalised-odds adjustment (as an example), as necessary.

9.2 Privacy

All the data sources are openly licensed and public: BTS data under the U.S. federal open data mandate, GHCN NOAA data as public domain and OpenFlights data under the

ODbL. No sensitive data requiring consent is gathered; only variables that are relevant in prediction are stored in the data lake, and none of the raw files with people's tail numbers is publicly published.

9.3 Security

All the following are in place: Shared repository access (GitHub branch-protection) via keys currently in the hands of all team members which are excluded in the .gitignore (API keys are stored as environment variables), MinIO role-based access control is used for access to the repositories, and all communication is by means of university-sanctioned Microsoft Teams/Outlook. The data lineage log is created and kept throughout the process, with every data transformation committed, following the approach of Stoudt et al. (2024).

10. References

Dai, M. (2024). A hybrid machine learning-based model for predicting flight delay through aviation big data. *Scientific Reports*, 14, Article 4603. <https://doi.org/10.1038/s41598-024-55217-z>

Faiza, & Khalil, K. (2023). Airline flight delays using artificial intelligence in COVID-19 with perspective analytics. *Journal of Intelligent & Fuzzy Systems*, 44(4), 6631–6653. <https://doi.org/10.3233/JIFS-222827>

Lambelho, M., Mitici, M., Pickup, S., & Marsden, A. (2020). Assessing strategic flight schedules at an airport using machine learning-based flight delay and cancellation predictions. *Journal of Air Transport Management*, 82, Article 101737. <https://doi.org/10.1016/j.jairtraman.2019.101737>

Li, Q., Wu, L., Guan, X., & Tian, Z. (2024). Interplay of network topologies in aviation delay propagation: A complex network and machine learning analysis. *Physica A: Statistical Mechanics and its Applications*, 638, Article 129622. <https://doi.org/10.1016/j.physa.2024.129622>

Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., & Galstyan, A. (2021). A survey on bias and fairness in machine learning. *ACM Computing Surveys*, 54(6), 1–35. <https://doi.org/10.1145/3457607>

Peterson, E. B., Neels, K., Barczi, N., & Graham, T. (2013). The economic cost of airline flight delay. *Journal of Transport Economics and Policy*, 47(1), 107–121. <https://doi.org/10.3828/jtep.2013.47.1.107>

Stoudt, S., Jernite, Y., Marshall, B., Marwick, B., Sharan, M., Whitaker, K., & Danchev, V. (2024). Ten simple rules for building and maintaining a responsible data science workflow. *PLOS Computational Biology*, 20(7), Article e1012232. <https://doi.org/10.1371/journal.pcbi.1012232>

Appendices

Appendix A: Data Dictionary

Table 5: Key fields by source dataset

Dataset	Field	Description	Type
BTS	FlightDate	Scheduled date of the flight	Date
BTS	Reporting_Airline	IATA carrier code	Categorical
BTS	Origin / Dest	Origin and destination airport IATA codes	Categorical
BTS	CRSDepTime / CRSArrTime	Scheduled local departure / arrival time	Time
BTS	DepDelayMinutes	Departure delay in minutes (0 if early/on-time)	Numeric
BTS	ArrDelayMinutes	Arrival delay in minutes — regression target	Numeric
BTS	Cancelled / CancellationCode	Cancellation flag and reason	Boolean / Categorical
BTS	CarrierDelay, WeatherDelay, NASDelay, SecurityDelay, LateAircraftDelay	Delay minutes attributed to each cause	Numeric
BTS	Distance	Route distance in miles	Numeric
NOAA GHCN-Daily	STATION, DATE	Weather station identifier and observation date	Categorical / Date
NOAA GHCN-Daily	PRCP, SNOW, TMAX, TMIN, AWND	Precipitation, snowfall, temperature extremes, wind speed	Numeric
OpenFlights	airport_id, IATA, ICAO	Airport identifiers used	Categorical

		to join to BTS Origin/Dest	
OpenFlights	latitude, longitude	Airport coordinates, used for nearest-station weather join	Numeric

Appendix B: Detailed Pipeline Process Design

Instead of giving the code for an actual implementation of each of the 5 "architecture" layers outlined in section 4, each layer is detailed at the design level, that is, what the layer is responsible for and what the layer passes on to the next.

- Lineage tracking of downloaded files is performed with retention of checksums, and with scheduled Airflow tasks pulling monthly BTS extractions over HTTPS, NOAA GHCN-Daily records over bulk FTP with API fallback, and finally picking up of the static OpenFlights CSV.
- Storage (MinIO): three buckets to store the data: the raw-zone (where the data lands before being immutable), curated-zone (where the data is stored after DuckDB cleanup and typed), and feature-zone (model-ready, versioned by training run, feature tables).
- Processing (DuckDB): SQL transformations deduplicate data on flight-leg keys, normalise the local-to-UTC conversion for departure and arrival time, join the flight information to the nearest station for the date of the flight, join the flight information to the airport data using IATA code, and impute missing delay-cause codes according to a set of rules.
- Temporal features (hour of day, day of week, holiday flag), carrier/airport delay-rate features (rolling the last 7 days) and features built from the graph (hub-centrality features based on Li et. al. (2024)) are added as features; models are trained on walk-forward windows and models are registered with their feature-set version for reproducibility.

Visualisation/automation: Visualisation is built into the Plotly Dash app, which is only fed by the feature-zone and a predictions table only updated by the nightly batch-scoring task scheduled by Airflow.

Appendix C: Model Development and Evaluation Plan

Table 6: Planned models and evaluation metrics.

Task	Candidate Models	Key Hyperparameters	Evaluation Metrics
Delay classification (delay / no delay)	Logistic Regression (baseline), Random Forest, XGBoost, LightGBM	Max depth, learning rate, n_estimators, class weight / scale_pos_weight	AUC-ROC, F1, precision/recall, Brier score, calibration curve
Delay-duration regression (minutes)	XGBoost Regressor, LightGBM Regressor	Max depth, learning rate, n_estimators, objective function	MAE, RMSE, R ²
Fairness check (per Mehrabi et al., 2021)	Same trained models, sliced by group	Grouping variables: carrier type, hub classification, route geography	Equalised-odds gap, demographic parity gap, per-group AUC/F1

Appendix D: Technology Justification Matrix

Table 7: Tool choices and alternatives considered

Tool Selected	Alternative(s) Considered	Rationale
DuckDB	Apache Spark	In-process SQL engine is sufficient for the ~2.5 GB Parquet feature store; avoids cluster setup and operational overhead
MinIO	AWS S3	Self-hosted and free during development; exposes an S3-compatible API, so migration to a cloud bucket later is low-risk
XGBoost /	CatBoost, neural	Strong reported performance on large tabular

LightGBM	networks	aviation data (Dai, 2024; Lambelho et al., 2020) with faster training and SHAP-based interpretability
Plotly Dash	Streamlit, Tableau	Full callback control for interactive filtering by carrier/airport/date, and remains open-source and Python-native
Apache Airflow	Prefect, cron	Mature DAG scheduling with built-in retry semantics and run history, needed for the nightly batch-scoring workflow